

Best build plan

Phase 1: MVP

Your first version should be a spreadsheet or simple database plus an AI prompt.

Input:

- a row of feedback text

Output:

- structured labels and summary

This can be built with:

- Google Sheets or Excel
- Claude or ChatGPT
- manual paste-in / copy-out workflow

This is enough to prove the idea.

Phase 2: Semi-automated workflow

After the prompt is working, automate the repetitive parts:

- new feedback goes into a sheet
- AI processes it
- output gets stored in a second sheet
- dashboard summarizes patterns

Phase 3: Dashboard and reporting

Create a simple view that shows:

- complaint trends
- sentiment trends
- category breakdowns
- urgent items
- sample quotes for each theme

That makes it look professional.

Best AI practices to follow

Use a strict output format

Ask the AI to return data in a fixed structure, ideally JSON or a table.

That reduces inconsistency and makes the results easier to analyze.

Separate extraction from interpretation

Have the AI do two passes if needed:

1. Extract facts from the feedback
2. Interpret business meaning

Example:

- Pass 1: “waited 40 minutes, food cold”
- Pass 2: “Service delay and food quality issue; high urgency”

This is more reliable than one giant prompt.

Keep categories limited

Do not create 30 categories. Start with 6–10 max.

Too many categories makes labeling messy and lowers accuracy.

Include a “miscellaneous” bucket

Some feedback will not fit neatly. That is normal.

Use human review for edge cases

Do not let AI fully auto-decide everything at first.

Review:

- unclear feedback
- mixed sentiment
- possible duplicates
- urgent complaints

Track confidence or uncertainty

Have the model mark items as:

- high confidence
- medium confidence
- low confidence

That helps you know when to review manually.

Protect privacy

If the data contains names, phone numbers, emails, order numbers, or other personal info:

- remove it

- mask it
- avoid using it unnecessarily

For a portfolio project, synthetic or anonymized data is best.

Version your prompts

Save each prompt version and note what changed.

That is a very professional habit and makes debugging much easier.

Measure quality

You need some way to know whether the system is good.

Simple metrics:

- category accuracy
 - consistency across repeated runs
 - how often human review changes the result
 - time saved vs manual processing
-

A very practical architecture

Use this simple flow:

Feedback source → **cleanup** → **AI classification** → **structured storage** → **dashboard** → **insights**

If you want to make it more concrete:

- source: reviews or survey comments
 - cleanup: remove duplicates, normalize text
 - AI step: tag category, sentiment, urgency
 - storage: spreadsheet / database
 - insights: charts and summary report
-

Efficient build strategy

The most efficient way is to avoid engineering too early.

Do not start with:

- custom app development
- complicated APIs
- fancy dashboards
- multi-source ingestion

- heavy automation

Start with:

- a spreadsheet
- a prompt
- a labeled sample set
- a simple summary report

Once that works, you can package it into something more polished.

What your final portfolio version should show

A strong portfolio version should demonstrate:

- you can turn messy unstructured text into structured business data
- you can identify recurring business problems
- you can prioritize issues by urgency and impact
- you can present insights clearly to decision-makers

That is professional value.

Suggested build deliverables

Your finished project could include:

1. a one-page project overview
2. sample dataset
3. labeling guide
4. prompt template
5. before/after examples
6. dashboard or report
7. short write-up on business impact

That is enough to look polished.

Best first version for you

If I were simplifying this as much as possible, I would do:

- 100 restaurant feedback comments
- 8 categories
- AI tags each comment

- spreadsheet stores results
- one dashboard shows top issues and praise
- one final summary report explains what management should do next

That is small enough to finish, but strong enough to show real value.